

Handoff: No Poker Face (Paper 1) — after calibration

6 October 2026 · Replaces the 5 October handoff · Upload with: this file, Paper1_No_Poker_Face_v0.2.pdf, No_Poker_Face_Calibration_and_Plan.pdf, no-poker-face-code.zip

Opening prompt: "Read this handoff, the calibration plan PDF and the v0.2 paper. Unzip the code package. In this chat we work only on No Poker Face. Read docs/findings-paper1.md and docs/cross-paper-ledger-entry.md in the zip. Next step is the one marked NEXT below."

1. Status

Item	State
Title	No Poker Face: Why AI Can't Keep the Answer to Itself
Paper	v0.2. Not yet revised with prior art or calibration
Prior art	Done. Gap confirmed: no work tests learner context → cue leakage in AI-written items. 30 references listed in the earlier chat and in the plan PDF's basis; 10 fields marked [verify]
Code	Built, mock-tested and run for real (zip). Resumable, cached, hard budget cap, spend ledger
Calibration	Done 2026-10-06. 640 items plus 160 secondary. \$24.28 spent of the \$65 approved
Main study	Not run. Awaiting approval: Batch API with an \$80 cap (recommended), or standard rates with a \$160 cap
Score	≈ 64/100 now → ≈ 82 done well (per-category table in the plan PDF, §6)

2. Calibration facts (exploratory; do not cite as findings)

- Invented domain (Tessellate 4), neutral condition: options-only accuracy 67.5% (chance 25%); full question 78.7%. Correct option longest in 71% of neutral items, pooled across both domains.
- H1 (W vs N): options-only −1.3 [−6.2, +3.7]. Invented-domain full question +10 [0, +21], but it reverses with rules on (WG vs NG −10). No consistent effect.
- Item-writing rules: keyed-longest 71% → 21%; invented-domain options-only still 54%.
- H2: notes wording copied into items in R (word-pair overlap 0.075 → 0.625 in the LLM API domain).
- Self-audit on request cut invented-domain full-question accuracy from 87.5% to 57.5% (n = 20). Pressure arm: no added compliance. Novelty: 3/40 near-repeats.
- Key check 99.4%. Judge κ: deciding fact 0.71, rule echo 0.64, dominant heuristic 0.64; three codes below 0.5 are descriptive only.

3. Decisions made, and why

- **Primary measure for the main study:** invented-domain full-question accuracy by cross-family solvers without the reference. It captures stem and option cues, and solver knowledge cannot contaminate it. Options-only is the key secondary measure. Primary contrast is W vs N; WG vs NG is a replication.
- **Size:** invented domain 40 slots × 8 conditions × 2 writers × 5 samples (≈94% power for 8 pts). LLM API domain 20 slots × 2 samples. Twenty new invented-domain slots still need writing.

- **Models:** writers Claude Opus 5.5 and GPT-5.6 Sol at defaults. Solvers only from the other family (Haiku 4.5 + Sonnet 5.5; GPT-5.6 Luna + Terra), with thinking off. Judges Sonnet 5.5 and GPT-5.6 Terra, blind. A third family needs a Google or OpenRouter key.
- **Framing kept:** sycophancy (helpfulness overriding the task), not deception. A human-written baseline, not chance, is the reference point. Fixed blueprint.
- **Prior art reframes the contributions:** heavy cues in LLM items are known (MedIWF, BenchMarker), so cite them. The contributions are H1, H2, pressure, self-audit and novelty. Ibrahim et al. (Nature 2026), Jain et al. (CHI 2026) and SafeTutors (2026) support the mechanism.

4. Cross-citation block — CHANGED 2026-10-06 (update the ledger and both handoffs)

- Paper 1 cites Paper 2 as: Suresh, V. (2026). Critic Ceiling: Why AI Review Runs Out of Real Problems. Study protocol, in preparation.
- Paper 2 cites Paper 1 as: Suresh, V. (2026). No Poker Face: Why AI Can't Keep the Answer to Itself. Working paper, in preparation.
- Shared-theme sentence, verbatim, once in each paper's Related work: "Both this paper and its companion examine a single theme from different sides: a model checking its own family's work misses what an independent check finds."
- **New:** Altitude Lock (doi:10.5281/zenodo.22847056) is cited in Paper 2, §3, **and in Paper 1**, Related work and Methods. The author decided this on 6 October.
- When either paper gets a DOI, swap it in here, in the other paper and in the ledger.

5. Rules for the next chat

- Never fabricate results. Mock-run outputs are synthetic and are never data. Exploratory calibration numbers are labelled as exploratory.
- Bulk runs go through the scripts only. Give a cost estimate and wait for approval. Keep the hard cap on: --budget.
- API keys are not stored in the zip. Supply them as environment variables. Rotate the keys pasted into the earlier chat.
- Record every AI error in findings-paper1.md, with who caught it.
- Deliver documents as PDF only. Never use real certification exam items.

6. Next steps, in order

#	Step
NEXT	Get approval for the main-run route. If Batch is chosen, add Batch API support to npf/clients.py for both providers and test it in mock mode
2	Write 20 more invented-domain slots, plus reference rules if needed. Add the compliance label to the pressure arm
3	Draft the pre-registration (hypotheses, primary measure, model, stopping rule, exclusions). The author posts it on OSF
4	Run the main study. Freeze the data. Run the pre-registered analysis. Independent review agent
5	Paper v1.0: new related work (30 refs, verify the 10 marked fields), methods, results, limitations, datasheet
6	Human baseline and coding (volunteers or paid); third family if a key arrives; Zenodo, then arXiv

7. Files

File	Contents
no-poker-face-code.zip	Code, data, notebook, calibration run (items, answers, codes, ledger, raw cached responses), docs
No_Poker_Face_Calibration_and_Plan.pdf	Results, design changes, costs, phases, scores
docs/findings-paper1.md	Commit to the paper's folder in the repo
docs/cross-paper-ledger-entry.md	Append to cross-paper-ledger.md in the repo root (no ledger file exists there yet)