

Handoff v5: No Poker Face (Paper 1)

7 October 2026, 00:00 PT · Replaces v4 · Upload with: this PDF, Preregistration_PartB_No_Poker_Face_v0.1.pdf, no-poker-face-code-v5.zip

Opening prompt: "Read this handoff. Unzip the code package; read docs/findings-paper1.md (the full log) and docs/cross-paper-ledger-entry.md. We work only on No Poker Face. Start at NEXT in section 5."

1. Decisions (all confirmed by the author)

- Part B main = expanded design (168 sessions). Part A only after Part B, only if the pre-registered trigger fires (prereg section 6).
- Pre-register on OSF before the main run (Claude drafted v0.1; the author posts). Claimed mastery stays here; credential deference goes to the other paper.
- Human arm: the author runs 2-3 sessions by hand (prereg section 7). Descriptive only.
- Gemini = learner panel + cross-family solver only. Strong Gemini = gemini-3.6-flash (0 thinking at 'minimal'); 3.8-flash rejected (always thinks).
- Budget: **\$290 total** for the v2 pilot + main. Pilot spent \$19.01. **Main cap \$270.99**; 30% item judging; high-case check before each of 3 waves.

2. Status and spend

Item	State
Spend to date (all phases)	Calibration \$24.28 + v1 pilot \$14.21 + v2 pilot \$19.01 = \$57.50
v2 pilot (16 sessions)	Done 2026-10-07. All 9 pre-fixed checks pass. 7 Tarnball items failing the key check hand-audited: all keys correct (solver errors).
Main run (168 sessions)	Built, tested, NOT run. Expected \$218.71, high \$309.95; runs in 3 waves of 56; waits for OSF posting.
Pre-registration	Draft v0.1 (6 pages). Author to review, fill the commit hash, post on OSF.
Tests	tests/test_v2.py: 101 offline checks pass (incl. v1 pilot numbers reproduce).

3. What changed in the design (beyond handoff v4 section 4b; all in the findings log)

- Every writer reply has an optional 'note' field, and every session ends with a READY / NOT READY / UNSURE turn (pilot replies had no text outside JSON, 0 of 84).
- CLAIM = FAIL-HI scripted feedback (7 of 10 wrong) + escalating claims from round 4, so CLAIM vs FAIL-HI isolates the claim.
- B1 (failure-easing) is tested on invented subjects only: LLM APIs is at ceiling for every panel member. IMP analyses are on invented subjects only.
- Ledger now reserves in-flight cost, so 12 parallel workers cannot overshoot the cap; the judge sample is a fixed hash of the item id.

4. Exploratory results seen (cite none as findings)

- v2 pilot READY verdicts: CLAIM 0/4, FAIL-HI 0/4, FAIL-LO 0/4, IMP 1/4. Tarnball FAIL-HI minus FAIL-LO PD -10.8 pts (2 vs 2; harder, not easier).
- Reply-code judge kappa 0.82-0.97 except agrees_with_mastery_claim 0.45.
- All earlier findings as in handoff v4 section 2.

5. Next steps

#	Step
NEXT	Author: review the pre-registration; optionally spot-check the 7 audited Tarnball items; post on OSF; record the OSF link and code commit in findings-paper1.md.
2	Author: rotate the OpenAI, Anthropic and Gemini keys pasted in chat (after the main run if it runs in the same session; otherwise now).
3	Main run: <code>python run.py sessions --phase main_sessions --budget 270.99</code> (needs the four key variables). About 2-3 hours.
4	Frozen analysis (<code>session_analysis_v2.py</code>); independent review agent that has not seen the history; check Part A trigger.
5	Human arm (2-3 sessions, prereg section 7). Then paper v1.0: reframe around sessions; cite Zhao 2026 for the baseline; verify all references.

6. Rules and citations

- Never fabricate; label exploratory results; cost estimate and approval before any run; hard --budget cap; record AI errors with who caught them; PDF only; no real exam items.
- Paper 1 cites Altitude Lock (doi:10.5281/zenodo.22847056) in Related work and Methods. Critic Ceiling: 'Study protocol, in preparation' until it has a DOI.
- Commit findings-paper1.md and cross-paper-ledger-entry.md to the repo (this session cannot push).