

No Poker Face, Part B: Addendum 1 to the pre-registration (osf.io/uya5q)

Gemini writer arm · Vijay Suresh · ORCID 0009-0004-1471-0561 · 7 October 2026

Summary. After the registration was posted, the author added a third writer family, Google Gemini, as a small extension: 28 sessions written by gemini-3.1-pro-preview. The registered main study (168 sessions; confirmatory tests B1 and B2) is unchanged. The arm is registered here before any of its data exist. It asks whether the main study's patterns replicate in a third writer family. No main-run outcome had been inspected when this decision was made.

1. What is added

Element	Gemini writer arm
Writer	gemini-3.1-pro-preview (Google), provider defaults (default thinking), one conversation per session
Design	4 subjects x 7 session types x 1 replicate = 28 sessions; same messages, same difficulty orders A/B (balanced 2/2 per type), same 7 rounds + readiness turn
Learner	gpt-5.6-luna (OpenAI), same gapped notes as the main study
Learner panel	claude-haiku-4-5-20251001 x 2 + gpt-5.6-luna, with the same three seeded gap profiles. No Gemini member, because no model evaluates its own family
Probe solvers	gpt-5.6-terra and claude-sonnet-5-5 (nohist, hist, spec, teach, teach_para, style-only)
Item solvers	Haiku 4.5, Sonnet 5.5, GPT-5.6 Luna, GPT-5.6 Terra; full question, options only, mismatched stem
Judges	claude-sonnet-5-5 and gpt-5.6-terra (item codes on a fixed 30% sample; reply codes on all non-empty replies)

2. Measures and analysis for the arm

- **Confirmatory tests are unchanged.** B1 and B2 (with Holm correction) are tested only on the registered 168 main-study sessions.
- In the arm, the B1 and B2 contrasts are reported as **replication estimates** (effect and 95% interval, same definitions), without a confirmatory test. With 3 vs 3 invented-subject sessions for B1 and 4 vs 4 for B2, they show direction only.
- PD in the arm = mean accuracy of the two Haiku panel members (the arm's equivalent of the two writer-invariant members). KF in the arm = mean of the Terra and Sonnet no-history probes. Because panel and solvers differ from the main study, arm values are compared with the main study for direction and pattern, not pooled into the confirmatory tests.
- All secondary measures of the registration (S1-S9) are computed for the arm with the same code and reported per writer; three-writer summaries are descriptive.
- Exclusions, missing-data rules and reporting of deviations follow the registration.

3. Data already seen

A 4-session feasibility pilot with this writer (Tarnball: C, FT, CLAIM, FAIL-HI; 7 October 2026; \$6.56) checked only feasibility: 280/280 valid questions, 4/4 readiness replies parsed, all notes non-empty, answer-key check 97.5%, cost \$1.64 per session. Its outcomes were not inspected. These 4 sessions are **not** part of the arm's sample; the arm is run fresh.

4. Code, budget and stopping

- Code: github.com/vjgits/research-papers, commit `b3f473b8293ad4451efbc45461922cc164c27fb2`, folder `No-Poker-Face-Why-AI-Cant-Keep-the-Answer-to-Itself/code-gemini-arm/no-poker-face` (phase 'gemini_arm'). The main run continues to use the frozen commit `a53ed958c872bfb392e756957bdf84163e588019`.
- Budget: hard cap \$60 (expected ~\$46 from the mini-pilot's measured cost). The run is a single wave; if the cap stops it early, the completed sessions are reported as such.
- Analysis of the arm happens together with the main analysis, after the main run is complete.