Recoverable Behavioral Failures: Multi-Model Experiment Pack v1.0

This pack operationalizes the paper's core claim as a runnable pilot benchmark. It is designed to test recoverable behavioral failure proxies across OpenAI, Anthropic, and Google Gemini models using deterministic scoring where possible. It does not report fabricated results; it provides the tasks, runner, scoring logic, and analysis plan needed to produce real results with API keys.

What can and cannot be run here

I cannot directly execute external multi-model API calls from this ChatGPT session because I do not have your OpenAI, Anthropic, or Gemini API keys, nor direct access to their paid endpoints from this container. The included runner is therefore a reproducible experiment harness. Once run with valid keys, the raw outputs and scored CSV can be inserted into the publishable paper as completed results.

Pilot task categories

Category	Failure mode tested	Deterministic signal
QI	Quantitative inconsistency	Expected numeric outputs are checked against prompt assumptions
PPG	Phantom placeholder persistence	Commands must not contain placeholder paths or template tokens
EBC	Editorial boundary contamination	Output must not include process notes, venue advice, scores, or chat-meta content
ACM	Audience/context miscalibration	Public-safe summary must avoid personal/sensitive workflow details
RETRY_PROXY	Retry/persistence strategy proxy	Model must propose alternative search strategy before declaring unknown

How to run

```
pip install -r requirements.txt
python run_eval.py   --models openai:gpt-5.5 anthropic:claude-sonnet-4-5 google:gemini-2.5-pro   --tasks task
python score_results.py   --raw results/raw_outputs.jsonl   --tasks tasks/evaluation_tasks.jsonl   --out resu
```

The model names above are placeholders and should be replaced with the currently enabled model IDs available to your API accounts. OpenAI documents its Python library and Responses API for direct model requests; Anthropic documents its Python SDK and Messages API; Google documents the Gemini Generate Content API and GenAI SDK.

Primary outcome measures

For each model and category, compute failure rate = failed tasks / total tasks. For multi-turn retry variants, compute first-pass success, after-retry success, and retry gap = after-retry success - first-pass success. For small pilot samples, report exact binomial or Wilson confidence intervals and avoid broad model-ranking claims.

Publication rule

Do not include the experiment as completed evidence in the paper until the raw outputs are generated and archived. Until then, describe it as a preregistered or proposed replication protocol. This protects the paper from the central methodological error the paper itself identifies: mixing process commentary and unfinished work into publishable claims.

Files included in the ZIP

README.md; requirements.txt; run_eval.py; score_results.py; tasks/evaluation_tasks.jsonl; empty results/ and paper_assets/ folders.